

---

# Homework 1

---

School of Information Science and Technology

Zhang Jun 12903136

March 15, 2019

## 1 SINGLE-RELATION QUERIES (36 PTS)

1. Consider the following relation:

Graph(n1 , n2)

A tuple (n1, n2) in Graph stores a directed edge from a node n1 to a node n2 in the corresponding graph. Your goal is to, for *every* node in the graph, count the number of outgoing edges of that node. Note that for nodes without any outgoing edges, their edge count would be zero; you need to output this as well. You can assume that:

- a) there are no duplicates or null values in the table; and
- b) every node in the graph is involved in at least one edge.

Answer:

```
select n1, count(n1)
from Graph
group by n1
union
select distinct n2, 0
from Graph
where n2 not in (
select n1
from Graph
)
```

2. Consider the following relation:

Trained(student, master, year)

A tuple (S, M, Y) in Trained specifies that a SQL Master M trained student S who graduated in year Y. Your goal is to find *the count of* SQL Masters who trained a student who graduated in the same year that 'Alice' or 'Bob' graduated.

Answer:

```
SELECT COUNT (*)
FROM (
  SELECT DISTINCT master
  FROM Trained
  WHERE year IN (
    SELECT DISTINCT year
    FROM Trained
    WHERE student = 'Alice' OR student = 'Bob') )
```

3. Consider the following relation:

DBMS(operator, system, performance)

A tuple (O, S, P) in DBMS specifies an operator O in system S and has the performance value P. Your goal is to find those systems whose operators achieves a higher performance value on average than the average performance value in a system named 'PostgreSQL'.

Answer:

```
SELECT system
FROM (SELECT system, AVG(performance) AS avg_per
FROM DBMS
GROUP BY system) AS new_DBMS
WHERE avg_per > (select AVG(performance)
FROM DBMS
WHERE system = 'PostgreSQL')
```

## 2 MULTI-RELATION QUERIES (40 PTS)

Consider the following relations representing student information at UIUC:

Mentorship(mentee\_sid, mentor\_sid)

Study(sid, credits)

Enrollment(did, sid)

Student(sid, street, city)

- A tuple (M1, M2) in Mentorship specifies that M2 is a mentor of another student M1.
- A tuple (S, C) in Study specifies that the student S has taken C credits.
- A tuple in Enrollment (D, S) specifies that student S is enrolled in department D.
- A (ST, S, C) in Student specifies that student ST lives on street S in city C.

1. Find all students who live in the same city and on the same street as their mentor.

Answer:

```
SELECT DISTINCT s1.sid
FROM Student s1, Student s2, Mentorship m
WHERE s1.city = s2.city
AND s1.street = s2.street
AND s1.sid = m.mentee_sid
AND s2.sid = m.mentor_sid
```

2. Find all students(i.e., distinct sid) who have taken more credits than the average credits of all of the students of their department.

Answer:

```
SELECT DISTINCT sid
FROM Study s, Enrollment e
WHERE E.sid = S.sid AND S.credits > (SELECT AVG(credits)
FROM Study s1, Enrollment e1
WHERE s1.sid = e1.sid AND
s.did = s1.did)
```

### 3 DATABASE MANIPULATION AND VIEWS (24 PTS)

1. In the Study relation, insert a new student, whose id is 66666 and has 0 credits.

Answer:

```
INSERT INTO Student VALUES (66666,NULL,NULL)
INSERT INTO Study VALUES (66666,0)
```

2. In the Study relation, delete students who have graduated (i.e., the ones who have more than 200 credits).

Answer:

```
DELETE
FROM Study
WHERE credits > 200
```

3. In the Study relation, add 2 credits for students who are mentors.

Answer:

```
UPDATE Study SET credits = credits + 2
WHERE sid in (
SELECT DISTINCT mentor_sid
FROM Mentorship )
```

4. Incoming students are those who have been accepted (i.e., exist in the **Student** relation) but have not registered in any department (i.e., do not exist in the **Enrollment** relation). Create a View that contains **sid** of all incoming students.

Answer:

```
CREATE VIEW incoming AS
SELECT Student.sid
FROM Student
WHERE Student.sid not in (
SELECT Enrollment.sid
FROM Enrollment)
```